

mk (02-09-2023)

```
In [ ]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [153]: 1 from sklearn.linear_model import LogisticRegression
          2 a=pd.read_csv(r"C:\USERS\user\Downloads\C4_framingham.csv")
          3 a
```

```
Out[153]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentH
0	1	39	4.0	0	0.0	0.0	0	
1	0	46	2.0	0	0.0	0.0	0	
2	1	48	1.0	1	20.0	0.0	0	
3	0	61	3.0	1	30.0	0.0	0	
4	0	46	3.0	1	23.0	0.0	0	
...	
4233	1	50	1.0	1	1.0	0.0	0	
4234	1	51	3.0	1	43.0	0.0	0	
4235	0	48	2.0	1	20.0	NaN	0	
4236	0	44	1.0	1	15.0	0.0	0	
4237	0	52	2.0	0	0.0	0.0	0	

```
In [154]: 1 from sklearn.linear_model import LogisticRegression
```

```
In [200]: 1 a=a.head(100)
          2 a
```

```
Out[200]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	di
0	1	39	4.0	0	0.0	0.0	0	0	
1	0	46	2.0	0	0.0	0.0	0	0	
2	1	48	1.0	1	20.0	0.0	0	0	
3	0	61	3.0	1	30.0	0.0	0	1	
4	0	46	3.0	1	23.0	0.0	0	0	
5	0	43	2.0	0	0.0	0.0	0	1	
6	0	63	1.0	0	0.0	0.0	0	0	
7	0	45	2.0	1	20.0	0.0	0	0	
8	1	52	1.0	0	0.0	0.0	0	1	
9	1	43	1.0	1	30.0	0.0	0	1	

```
In [201]: 1 a.columns
```

```
Out[201]: Index(['male', 'age', 'education', 'currentSmoker', 'cigsPerDay', 'BPMeds',
                  'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP',
                  'diaBP', 'BMI', 'heartRate', 'glucose', 'TenYearCHD'],
                  dtype='object')
```

```
In [202]: 1 b=a[['male', 'age', 'education', 'currentSmoker', 'cigsPerDay', 'BPMeds',
                'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP',
                'diaBP', 'BMI', 'heartRate', 'glucose', 'TenYearCHD']]
          2
          3
          4 b
```

```
Out[202]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	di
0	1	39	4.0	0	0.0	0.0	0	0	
1	0	46	2.0	0	0.0	0.0	0	0	
2	1	48	1.0	1	20.0	0.0	0	0	
3	0	61	3.0	1	30.0	0.0	0	1	
4	0	46	3.0	1	23.0	0.0	0	0	
5	0	43	2.0	0	0.0	0.0	0	1	
6	0	63	1.0	0	0.0	0.0	0	0	
7	0	45	2.0	1	20.0	0.0	0	0	
8	1	52	1.0	0	0.0	0.0	0	1	
9	1	43	1.0	1	30.0	0.0	0	1	

```
In [203]: 1 c=b.iloc[:,0:15]
          2 d=b.iloc[:, -1]
```

```
In [204]: 1 c.shape
```

```
Out[204]: (10, 15)
```

```
In [205]: 1 d.shape
```

```
Out[205]: (10,)
```

```
In [ ]: 1
```

```
In [206]: 1 from sklearn.preprocessing import StandardScaler
          2 fs=StandardScaler().fit_transform(c)
          3 fs
```

```
Out[206]: array([[ 1.22474487, -1.28931674,  2.          , -1.          , -0.96754461,
                   0.          ,  0.          , -0.81649658,  0.          , -1.41062997,
                   -1.27734618, -1.15306967,  0.05752806,  0.14484136, -0.70928138],
                 [-0.81649658, -0.34918995,  0.          , -1.          , -0.96754461,
                   0.          ,  0.          , -0.81649658,  0.          ,  0.20235648,
                   -0.63004237, -0.35029965,  0.58486866,  1.59325501, -0.8106073 ],
                 [ 1.22474487, -0.0805823 , -1.          ,  1.          ,  0.60569866,
                   0.          ,  0.          , -0.81649658,  0.          ,  0.05572135,
                   -0.34954405, -0.42327874, -0.43086123, -0.33796318, -1.41856277],
                 [-0.81649658,  1.66536746,  1.          ,  1.          ,  1.39232029,
                   0.          ,  0.          ,  1.22474487,  0.          , -0.53081918,
                   0.62141165,  0.67140766,  0.53992486, -1.30357228,  1.92519233],
                 [-0.81649658, -0.34918995,  1.          ,  1.          ,  0.84168515,
                   0.          ,  0.          , -0.81649658,  0.          ,  1.22880241,
                   -0.24166009, -0.13136237, -1.10202199,  0.62764591,  0.10132591],
                 [-0.81649658, -0.75210143,  0.          , -1.          , -0.96754461,
                   0.          ,  0.          ,  1.22474487,  0.          , -0.4428381 ,
                   1.91601926,  1.76609405,  1.05528044, -0.14484136,  1.51988868],
                 [-0.81649658,  1.93397511, -1.          , -1.          , -0.96754461,
                   0.          ,  0.          , -0.81649658,  0.          , -1.11735971,
                   0.10356861, -1.08009058,  1.89722763, -1.78637683,  0.10132591],
                 [-0.81649658, -0.48349378,  0.          ,  1.          ,  0.60569866,
                   0.          ,  0.          , -0.81649658,  0.          ,  2.04995915,
                   -1.5362677 , -1.08009058, -1.52748997,  0.04828045, -0.60795547],
                 [ 1.22474487,  0.45663301, -1.          , -1.          , -0.96754461,
                   0.          ,  0.          ,  1.22474487,  0.          ,  0.49562675,
                   0.25460616,  0.2335331 , -0.12524339, -0.24140227, -0.50662956],
                 [ 1.22474487, -0.75210143, -1.          ,  1.          ,  1.39232029,
                   0.          ,  0.          ,  1.22474487,  0.          , -0.53081918,
                   1.1392547 ,  1.54715677, -0.94921307,  1.40013319,  0.40530365]])
```

```
In [207]: 1 logr=LogisticRegression()
          2 logr.fit(fs,d)
```

```
Out[207]: LogisticRegression()
```

```
In [208]: 1 e=[[2,5,77,8,5,2.3,5.2,1,1.2,16,56,52,45,25,65]]
```

```
In [209]: 1 prediction=logr.predict(e)
          2 prediction
```

```
Out[209]: array([1], dtype=int64)
```

```
In [210]: 1 logr.classes_
```

```
Out[210]: array([0, 1], dtype=int64)
```

```
In [211]: 1 logr.predict_proba(e)[0][0]
```

```
Out[211]: 1.9817379515174594e-08
```

```
In [212]: 1 import re
          2 from sklearn.datasets import load_digits
          3 import numpy as np
          4 import pandas as pd
          5 import matplotlib.pyplot as plt
          6 import seaborn as sns
```

```
In [213]: 1 from sklearn.linear_model import LogisticRegression
          2 from sklearn.model_selection import train_test_split
```

```
In [214]: 1 digits=load_digits()
          2 digits
```


```
[[ 0.,  0.,  0., ...,  0.,  0.,  0.],
 [ 0.,  8., 16., ..., 16.,  8.,  0.],
 [ 0.,  1.,  8., ..., 12.,  1.,  0.]...],
'DESCR': ".. _digits_dataset:\n\nOptical recognition of handwritten digits
dataset\n-----\n\n**Data Set Characteristics:**\n\n    :Number of Instances: 1797\n    :Number of Attributes: 64\n    :Attribute Information: 8x8 image of integer pixels in the range 0..16.\n    :Missing Attribute Values: None\n    :Creator: E. Alpaydin (alpaydin '@' boun.edu.tr)\n    :Date: July; 1998\n\nThis is a copy of the test set of the UCI ML hand-written digits datasets\nhttps://archive.ics.uci.edu/ml/datasets/Optical+Recognition+of+Handwritten+Digits\n\nThe data set contains images of hand-written digits: 10 classes where each class refers to a digit.\n\nPreprocessing programs made available by NIST were used to extract\nnormalized bitmaps of handwritten digits from a preprinted form. From a total of 43 people, 30 contributed to the training set and different 13 to the test set. 32x32 bitmaps are divided into nonoverlapping blocks of 4x4 and the number of on pixels are counted in each block. This generates an input matrix of 8x8 where each element is an integer in the range 0..16. This reduces dimensionality and gives invariance to small distortions.\n\nFor info on NIST preprocessing routines, see M. D. Garris, J. L. Blue, G. A. Cottrell, D. L. Birmaher, J. Geist, D. J. Graham, G. A. Jones, and
```

```
In [215]: 1 plt.figure(figsize=(20,4))
```

```
Out[215]: <Figure size 1440x288 with 0 Axes>
```

```
<Figure size 1440x288 with 0 Axes>
```

```
In [216]: 1 for index,(image,label) in enumerate(zip(digits.data[0:5],digits.target[0:
2         plt.subplot(1,5,index+1)
3         plt.imshow(np.reshape(image,(8
4         ,8)),cmap=plt.cm.gray)
5         plt.title('Number:%i\n'%label,fontsize=4)
```



```
In [217]: 1 x_train,x_test,y_train,y_test=train_test_split(digits.data,digits.target,t
```

```
In [218]: 1 print(x_train.shape)
2 print(x_test.shape)
3 print(y_train.shape)
4 print(y_test.shape)
```

```
(1257, 64)
(540, 64)
(1257,)
(540,)
```

```
In [219]: 1 logre=LogisticRegression(max_iter=10000)
2 logre.fit(x_train,y_train)
3
```

```
Out[219]: LogisticRegression(max_iter=10000)
```

```
In [220]: 1 print(logre.predict(x_test))
```

```
[6 6 9 4 3 3 9 1 3 6 0 6 5 9 8 7 0 8 3 7 1 0 6 4 4 8 4 1 9 4 1 6 9 8 3 4 9
 1 4 3 1 9 0 1 6 2 6 5 3 0 0 1 5 2 7 4 3 1 2 7 0 1 4 0 1 5 5 7 1 9 7 9 2 9
 2 6 7 4 2 5 8 9 1 5 5 2 1 3 1 3 7 2 7 5 3 9 0 9 3 0 4 0 1 2 5 9 6 9 6 7 8
 4 5 6 9 8 9 9 2 5 7 1 2 4 8 8 3 2 9 5 1 2 4 5 1 9 3 2 2 9 9 1 3 5 0 7 3 4
 9 7 8 3 5 9 4 4 8 5 5 0 0 7 9 9 2 8 5 5 2 7 1 3 7 5 7 6 5 0 6 6 8 9 5 5 6
 5 1 8 4 2 3 6 3 4 1 7 4 2 3 3 6 2 2 4 3 6 1 7 2 1 9 8 6 6 5 0 1 0 5 6 2 1
 4 1 6 1 8 3 6 5 9 4 3 3 9 3 1 5 1 4 8 9 6 1 3 9 9 7 0 9 8 6 9 5 8 9 7 7 2
 3 4 1 9 8 6 9 1 3 6 4 1 4 4 9 1 2 7 6 7 3 8 8 9 2 7 6 7 9 9 1 0 7 3 0 2 6
 7 0 5 1 5 3 2 0 3 7 3 5 3 0 4 8 2 8 8 1 2 9 6 3 3 0 7 5 8 5 2 0 8 3 1 5 4
 3 6 0 7 7 4 4 1 5 3 7 1 8 7 7 7 6 4 1 5 1 8 0 9 8 8 2 1 6 1 9 2 2 1 9 1 7
 7 6 4 7 9 2 9 4 0 3 2 4 9 2 4 7 5 9 4 2 3 8 2 1 0 3 3 8 6 5 5 1 7 8 2 4 9
 8 3 4 1 8 0 1 7 7 1 5 8 1 5 4 1 1 8 7 9 6 6 0 2 8 5 2 0 3 7 2 4 8 4 3 2 9
 0 6 6 3 8 1 2 0 1 4 9 9 0 6 1 4 0 7 7 5 6 3 2 9 7 2 9 1 9 1 5 0 9 1 2 6 8
 5 6 0 7 0 7 8 7 2 4 3 7 7 9 1 7 7 5 6 4 2 9 0 5 5 7 7 0 6 4 3 0 4 4 5 0 9
 5 0 4 6 5 0 2 9 8 9 8 5 5 0 7 0 1 5 0 0 2 8]
```

```
In [221]: 1 import numpy as np
          2 import pandas as pd
          3 import matplotlib.pyplot as plt
          4 import seaborn as sns
```

```
In [270]: 1 a=pd.read_csv(r"C:\USERS\user\Downloads\C4_framingham.csv")
```

```
In [271]: 1 a['male'].value_counts()
```

```
Out[271]: 0    2419
          1    1819
          Name: male, dtype: int64
```

```
In [272]: 1 x=b.drop('male',axis=1)
          2 y=b['male']
          3 print(b)
```

```
-----
KeyError                                Traceback (most recent call last)
<ipython-input-272-1b5f8ccb0541> in <module>
----> 1 x=b.drop('male',axis=1)
      2 y=b['male']
      3 print(b)

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\frame.py in drop(self,
labels, axis, index, columns, level, inplace, errors)
    4306         weight 1.0      0.8
    4307         """
-> 4308         return super().drop(
    4309             labels=labels,
    4310             axis=axis,

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\generic.py in drop(self,
labels, axis, index, columns, level, inplace, errors)
    4151         for axis, labels in axes.items():
    4152             if labels is not None:
-> 4153                 obj = obj._drop_axis(labels, axis, level=level, error
s=errors)
    4154
    4155         if inplace:

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\generic.py in _drop_axis
(self, labels, axis, level, errors)
    4186         new_axis = axis.drop(labels, level=level, errors=errors)
    4187         else:
-> 4188             new_axis = axis.drop(labels, errors=errors)
    4189             result = self.reindex(**{axis_name: new_axis})
    4190

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexes\base.py in drop
p(self, labels, errors)
    5589         if mask.any():
    5590             if errors != "ignore":
-> 5591                 raise KeyError(f"{labels[mask]} not found in axis")
    5592             indexer = indexer[~mask]
    5593         return self.delete(indexer)

KeyError: '['male'] not found in axis"
```

In [273]:

```
1 g1={"Mention Count":{"Mention Count":1,'b':2}}
2 df=df.replace(g1)
3 print(df)
```


	User ID	Username \
0	132131	flong
1	289683	hinesstephanie
2	779715	robertttran
3	696168	pmason
4	704441	noah87
...
49995	491196	uberg
49996	739297	jessicamunoz
49997	674475	lynnccunningham
49998	167081	richardthompson
49999	311204	daniel29

	Tweet	Retweet Count \
0	Station activity person against natural majori...	85
1	Authority research natural life material staff...	55
2	Manage whose quickly especially foot none to g...	6
3	Just cover eight opportunity strong policy which.	54
4	Animal sign six data good or.	26
...
49995	Want but put card direction know miss former h...	64
49996	Provide whole maybe agree church respond most ...	18
49997	Bring different everyone international capital...	43
49998	Than about single generation itself seek sell ...	45
49999	Here morning class various room human true bec...	91

	Mention Count	Follower Count	Verified	Bot Label	Location
\					
0	1	2353	False	1	Adkinston
1	5	9617	True	0	Sanderston
2	2	4363	True	0	Harrisonfurt
3	5	2242	True	1	Martinezberg
4	3	8438	False	1	Camachoville
...
49995	0	9911	True	1	Lake Kimberlyburgh
49996	5	9900	False	1	Greenbury
49997	3	6313	True	1	Deborahfort
49998	1	6343	False	0	Stephenside
49999	4	4006	False	0	Novakberg

	Created At	Hashtags
0	2020-05-11 15:29:50	NaN
1	2022-11-26 05:18:10	both live
2	2022-08-08 03:16:54	phone ahead
3	2021-08-14 22:27:05	ever quickly new I
4	2020-04-13 21:24:21	foreign mention
...
49995	2023-04-20 11:06:26	teach quality ten education any
49996	2022-10-18 03:57:35	add walk among believe
49997	2020-07-08 03:54:08	onto admit artist first
49998	2022-03-22 12:13:44	star
49999	2022-12-03 06:11:07	home

[50000 rows x 11 columns]

```
In [274]: 1 from sklearn.model_selection import train_test_split
          2 x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [275]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [276]: 1 rfc=RandomForestClassifier()
          2 rfc.fit(x_train,y_train)
```

Out[276]: RandomForestClassifier()

```
In [277]: 1 parameters={'max_depth':[1,2,3,4,5],
          2             'min_samples_leaf':[5,10,15,20,25],
          3             'n_estimators':[10,20,30,40,50]}
```

```
In [278]: 1 from sklearn.model_selection import GridSearchCV
```

```
In [279]: 1 grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring=
          2 grid_search.fit(x_train,y_train)
```

Out[279]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
 'min_samples_leaf': [5, 10, 15, 20, 25],
 'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')

```
In [280]: 1 grid_search.best_score_
```

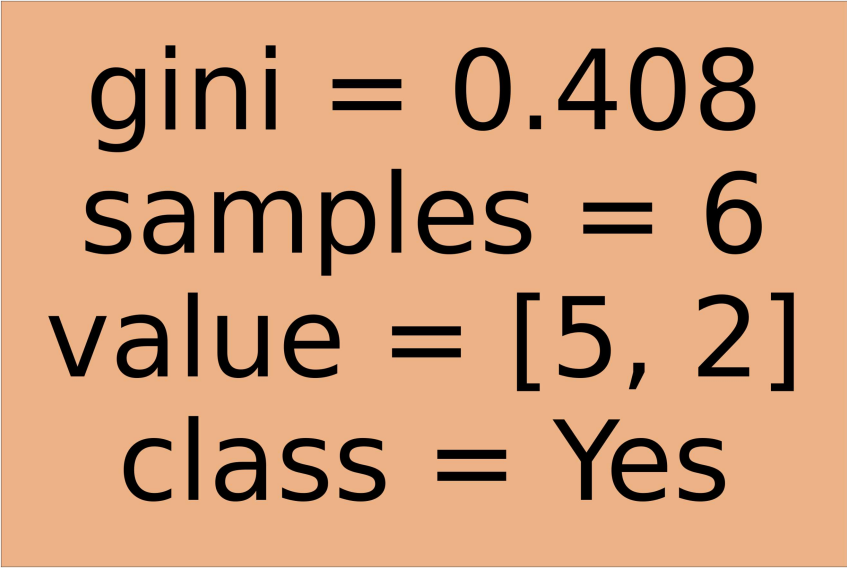
Out[280]: 0.5833333333333333

```
In [281]: 1 rfc_best=grid_search.best_estimator_
```

```
In [282]: 1 from sklearn.tree import plot_tree
```

```
In [283]: 1 plt.figure(figsize=(80,40))
          2 plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes',
          3                                     'No'])
```

```
Out[283]: [Text(2232.0, 1087.2, 'gini = 0.408\nsamples = 6\nvalue = [5, 2]\nclass = Yes')]
          s')]
```



gini = 0.408
samples = 6
value = [5, 2]
class = Yes

```
In [ ]:
```

```
1
```