

mk 02-09-2023

```
In [ ]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [238]: 1 from sklearn.linear_model import LogisticRegression
          2 a=pd.read_csv(r"C:\USERS\user\Downloads\C5_health care diabetes.csv")
          3 a
```

Out[238]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.288
...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0.201
764	2	122	70	27	0	36.8	0.248
765	5	121	72	23	112	26.2	0.134
766	1	126	60	0	0	30.1	0.158
767	1	93	70	31	0	30.4	0.232

```
In [303]: 1 a=a.head(10)
          2 a
```

Out[303]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.288
5	5	116	74	0	0	25.6	0.201
6	3	78	50	32	88	31.0	0.248
7	10	115	0	0	0	35.3	0.134
8	2	197	70	45	543	30.5	0.158
9	8	125	96	0	0	0.0	0.232

```
In [304]: 1 from sklearn.linear_model import LogisticRegression
```

```
In [305]: 1 a.columns
```

```
Out[305]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
                'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
                dtype='object')
```

```
In [306]: 1 b=a[['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
                2         'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome']]
           3 b
```

```
Out[306]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.288
5	5	116	74	0	0	25.6	0.201
6	3	78	50	32	88	31.0	0.248
7	10	115	0	0	0	35.3	0.134
8	2	197	70	45	543	30.5	0.158
9	8	125	96	0	0	0.0	0.232

```
In [307]: 1 c=b.iloc[:,0:15]
           2 d=b.iloc[:, -1]
```

```
In [308]: 1 c.shape
```

```
Out[308]: (10, 9)
```

```
In [309]: 1 d.shape
```

```
Out[309]: (10,)
```

```
In [ ]: 1
```

```
In [310]: 1 from sklearn.preprocessing import StandardScaler
          2 fs=StandardScaler().fit_transform(c)
          3 fs
```

```
Out[310]: array([[ 0.48154341,  0.54447108,  0.50126452,  0.88534503, -0.55391299,
                   0.55143407,  0.19208855,  1.25171488,  0.81649658],
                  [-1.02327975, -1.11261482,  0.25474099,  0.5335523 , -0.55391299,
                   -0.10392051, -0.25268024, -0.43499311, -1.22474487],
                  [ 1.08347268,  1.46507436,  0.17256648, -1.16677921, -0.55391299,
                   -0.41287339,  0.2646052 , -0.34621901,  0.81649658],
                  [-1.02327975, -1.00740302,  0.25474099,  0.18175957,  0.02915332,
                   0.03651261, -0.54919276, -1.32273416, -1.22474487],
                  [-1.32424438,  0.25513862, -0.81352767,  0.88534503,  0.48816296,
                   1.44084387,  2.86875866, -0.2574449 ,  0.81649658],
                  [ 0.18057878, -0.29722334,  0.58343903, -1.16677921, -0.55391299,
                   -0.1975426 , -0.4944024 , -0.52376722, -1.22474487],
                  [-0.42135049, -1.29673547, -0.40265511,  0.70944866, -0.00806368,
                   0.30801666, -0.41866279, -0.87886364,  0.81649658],
                  [ 1.68540194, -0.32352629, -2.45701791, -1.16677921, -0.55391299,
                   0.71059162, -0.60237164, -0.61254132, -1.22474487],
                  [-0.72231512,  1.83331567,  0.41909001,  1.47166623,  2.81422536,
                   0.26120561, -0.56369609,  1.51803719,  0.81649658],
                  [ 1.08347268, -0.06049679,  1.48735867, -1.16677921, -0.55391299,
                   -2.59426794, -0.44444649,  1.6068113 ,  0.81649658]])
```

```
In [311]: 1 logr=LogisticRegression()
          2 logr.fit(fs,d)
```

```
Out[311]: LogisticRegression()
```

```
In [312]: 1 e=[[2,5,77,8,56,52,45,25,65]]
```

```
In [313]: 1 prediction=logr.predict(e)
          2 prediction
```

```
Out[313]: array([1], dtype=int64)
```

```
In [314]: 1 logr.classes_
```

```
Out[314]: array([0, 1], dtype=int64)
```

```
In [315]: 1 logr.predict_proba(e)[0][0]
```

```
Out[315]: 0.0
```

```
In [316]: 1 import re
          2 from sklearn.datasets import load_digits
          3 import numpy as np
          4 import pandas as pd
          5 import matplotlib.pyplot as plt
          6 import seaborn as sns
```

```
In [317]: 1 from sklearn.linear_model import LogisticRegression
          2 from sklearn.model_selection import train_test_split
```

```
In [318]: 1 digits=load_digits()
          2 digits


          3 [ 0.,  4., 10., ..., 10.,  0.,  0.],
          4 [ 0.,  8., 16., ..., 16.,  8.,  0.],
          5 [ 0.,  1.,  8., ..., 12.,  1.,  0.]])],
          6
          7 'DESCR': ".. _digits_dataset:\n\nOptical recognition of handwritten digits
          8 dataset\n-----\n\n**Data Set C
          9 characteristics:**\n\n      :Number of Instances: 1797\n      :Number of Attribu
         10 tes: 64\n      :Attribute Information: 8x8 image of integer pixels in the ran
         11 ge 0..16.\n      :Missing Attribute Values: None\n      :Creator: E. Alpaydin
         12 (alpaydin '@' boun.edu.tr)\n      :Date: July; 1998\n\nThis is a copy of the
         13 test set of the UCI ML hand-written digits datasets\nhttps://archive.ics.uc
         14 i.edu/ml/datasets/Optical+Recognition+of+Handwritten+Digits\n\nThe data set
         15 contains images of hand-written digits: 10 classes where\neach class refers
         16 to a digit.\n\nPreprocessing programs made available by NIST were used to e
         17 xtract\nnormalized bitmaps of handwritten digits from a preprinted form. Fr
         18 om a\ntotal of 43 people, 30 contributed to the training set and different
         19 13\nto the test set. 32x32 bitmaps are divided into nonoverlapping blocks o
         20 f\n4x4 and the number of on pixels are counted in each block. This generate
         21 s\nan input matrix of 8x8 where each element is an integer in the range\n
         22 0..16. This reduces dimensionality and gives invariance to small\ndistortio
         23 ns.\n\nFor info on NIST preprocessing routines, see M. D. Garris, J. L. Blu
         24 ck, T. S. Huang, D. J. D'Angelo, J. F. Flourens, D. J. Frazer, C. A. Gottlieb,
         25 J. H. J. Leung, P. Y. Yuen, and J. C. Zide, NIST Special Publication 900-1, NIST
         26 Technology Series, Gaithersburg, MD, 1997.
```

```
In [319]: 1 plt.figure(figsize=(20,4))
```

Out[319]: <Figure size 1440x288 with 0 Axes>

<Figure size 1440x288 with 0 Axes>

```
In [320]: 1 for index,(image,label) in enumerate(zip(digits.data[0:5],digits.target[0:
          2         plt.subplot(1,8,index+1)
          3         plt.imshow(np.reshape(image,(8
          4         ,8)),cmap=plt.cm.gray)
          5         plt.title('Number:%i\n'%label,fontsize=10)
```



```
In [321]: 1 x_train,x_test,y_train,y_test=train_test_split(digits.data,digits.target,t
```

```
In [322]: 1 print(x_train.shape)
          2 print(x_test.shape)
          3 print(y_train.shape)
          4 print(y_test.shape)
```

```
(1257, 64)
(540, 64)
(1257,)
(540,)
```

```
In [323]: 1 logre=LogisticRegression(max_iter=10000)
          2 logre.fit(x_train,y_train)
          3
```

```
Out[323]: LogisticRegression(max_iter=10000)
```

```
In [324]: 1 print(logre.predict(x_test))
```

```
[6 1 1 1 0 7 7 2 1 7 4 7 5 1 4 4 4 3 5 5 5 8 4 7 6 8 6 2 9 0 7 9 0 5 8 7 5
 1 8 9 0 2 7 3 3 1 8 1 9 4 4 5 6 5 2 8 3 4 1 4 3 0 5 4 0 2 8 8 0 9 5 4 7 3
 5 1 3 5 8 8 5 3 1 4 0 1 8 4 8 0 7 4 4 6 7 0 9 9 4 1 0 2 7 5 0 5 8 2 0 6 5
 8 2 5 1 9 4 3 0 1 1 1 6 4 3 8 1 5 0 6 9 4 0 2 2 2 6 7 8 2 7 1 2 9 9 4 1 8
 0 6 6 4 2 2 3 5 4 1 2 4 6 2 9 1 0 0 5 4 1 8 4 5 6 1 1 5 7 8 7 3 0 2 1 8 6
 2 8 1 4 1 8 1 3 0 3 4 8 3 4 7 4 4 4 7 3 0 9 5 7 2 1 5 3 2 3 3 0 0 9 5 0 8
 1 6 9 4 8 8 5 8 7 8 2 5 9 5 9 2 5 5 4 3 9 0 5 5 6 1 2 1 2 1 1 4 6 8 1 1 9
 0 2 5 3 7 3 2 3 2 7 6 6 3 7 4 3 9 9 9 6 1 1 0 2 9 6 9 8 4 9 0 0 9 3 0 6 9
 6 5 6 7 2 7 6 0 2 5 0 8 9 0 9 3 1 1 4 3 7 6 3 9 2 3 1 8 1 7 4 6 3 2 1 6 2
 8 3 1 9 4 3 7 9 6 4 2 2 2 8 2 9 1 4 7 6 2 6 2 7 7 0 5 5 2 2 3 4 5 3 1 7 7
 0 5 1 6 1 5 2 3 6 2 4 1 2 4 5 1 3 1 5 5 4 8 4 1 3 0 1 5 7 8 3 3 7 5 6 5 1
 5 8 2 0 8 3 1 6 3 2 2 2 9 4 0 2 4 6 3 7 6 8 6 6 1 9 9 8 1 1 2 1 5 1 4 9 5
 7 0 1 6 8 0 9 8 5 6 3 3 4 5 0 2 8 2 9 0 6 9 9 2 7 0 8 8 3 6 6 9 4 8 2 6 3
 1 8 9 5 5 6 0 4 2 8 4 8 9 0 0 7 1 3 9 0 3 4 7 5 3 6 0 2 3 1 7 9 7 3 3 8 9
 0 7 1 0 8 8 7 2 1 2 4 1 3 8 4 3 7 2 8 0 9 7]
```

```
In [325]: 1 import numpy as np
          2 import pandas as pd
          3 import matplotlib.pyplot as plt
          4 import seaborn as sns
```

```
In [326]: 1 a=pd.read_csv(r"C:\USERS\user\Downloads\C5_health care diabetes.csv")
```

```
In [329]: 1 a=a.head(10)
          2 a
```

Out[329]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.288
5	5	116	74	0	0	25.6	0.201
6	3	78	50	32	88	31.0	0.248
7	10	115	0	0	0	35.3	0.134
8	2	197	70	45	543	30.5	0.158
9	8	125	96	0	0	0.0	0.232

```
In [343]: 1 a['Outcome'].value_counts()
```

Out[343]: 1 6  
0 4  
Name: Outcome, dtype: int64

```
In [358]: 1 x=b.drop('Outcome',axis=1)
          2 y=b['Outcome']
          3 print(b)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
5	5	116	74	0	0	25.6	
6	3	78	50	32	88	31.0	
7	10	115	0	0	0	35.3	
8	2	197	70	45	543	30.5	
9	8	125	96	0	0	0.0	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
5	0.201	30	0
6	0.248	26	1
7	0.134	29	0
8	0.158	53	1
9	0.232	54	1

```
In [359]: 1 g1={"Outcome":{'Outcome':1,'b':2}}
          2 a=a.replace(g1)
          3 print(a)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
5	5	116	74	0	0	25.6	
6	3	78	50	32	88	31.0	
7	10	115	0	0	0	35.3	
8	2	197	70	45	543	30.5	
9	8	125	96	0	0	0.0	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
5	0.201	30	0
6	0.248	26	1
7	0.134	29	0
8	0.158	53	1
9	0.232	54	1

```
In [360]: 1 from sklearn.model_selection import train_test_split
          2 x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [361]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [362]: 1 rfc=RandomForestClassifier()
          2 rfc.fit(x_train,y_train)
```

Out[362]: RandomForestClassifier()

```
In [363]: 1 parameters={'max_depth':[1,2,3,4,5],
          2               'min_samples_leaf':[5,10,15,20,25],
          3               'n_estimators':[10,20,30,40,50]}
```

```
In [364]: 1 from sklearn.model_selection import GridSearchCV
```

```
In [365]: 1 grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring=
          2 grid_search.fit(x_train,y_train)
```

Out[365]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
 param\_grid={'max\_depth': [1, 2, 3, 4, 5],  
 'min\_samples\_leaf': [5, 10, 15, 20, 25],  
 'n\_estimators': [10, 20, 30, 40, 50]},  
 scoring='accuracy')

```
In [366]: 1 grid_search.best_score_
```

```
Out[366]: 0.5833333333333333
```

```
In [367]: 1 rfc_best=grid_search.best_estimator_
```

```
In [368]: 1 from sklearn.tree import plot_tree
```

```
In [369]: 1 plt.figure(figsize=(80,40))  
2 plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Ye  
3
```

```
Out[369]: [Text(2232.0, 1087.2, 'gini = 0.49\nsamples = 6\nvalue = [3, 4]\nnclass = N  
o')]
```

**gini = 0.49  
samples = 6  
value = [3, 4]  
class = No**

```
In [ ]: 1
```